

## Computational Neuroscience

### Course Highlights:

- Some light neurobiology
- PCA and eigenbases
- Backpropagation
- Circuit analysis for neuromodels
- Eigenfaces

### WEEK 1 – Introduction to Computational Neuroscience

#### 1.1 Course Introduction

- Descriptive Models
  - how do neurons respond to stimuli and how is that quantitatively encoded
  - how can we extract info from neurons (**decoding**)
- How can we simulate a single neuron?
- Why do brain circuits operate the way they do?

#### *At the end of the course...*

- should be able to quantitatively describe what is going on with a neuron or a network
- simulate behavior of neurons
- formulate computational neurons

#### 1.2 Descriptive Models

- Goal: explain how brains generate behaviors
- Going to characterize what nervous systems do, how they function, and why they operate in particular ways
  - Descriptive models (what)
  - Mechanistic models (how the neural system does what it does)
  - Interpretative models (why)
- Output from brain cell → *action potential*
- **Def: receptive field:**
  - **Specific properties of a sensory stimulus that generate a strong response from the cell**
- Retina – layer of tissue at the back of the eyes
  - Inverted image projected onto back of the eyes
  - Retinal ganglion cells – conveying information about the image to other parts of the brain
- Information from the retina passed to the *Lateral Geniculate Nucleus (LGN)* which then passes information to the Primary Visual Cortex V1.

- Center surround LGN receptive fields are displaced because of the preferred orientation of the primary visual cortex

### 1.3 Mechanistic and Interpretive Models


- *Efficient coding hypothesis* – suppose goal is to represent images as faithfully as possible using neurons with receptive fields
- Given image  $I$ , we can reconstruct with a linear combination of receptive fields multiplied by the respective neural response
- We care about **minimizing** the total square pixel wise error and also making sure they're as independent as possible?
- Idea is like start with random receptive field and then run the coding algorithm on natural image patches
  - What is the efficient coding algorithm?
    - Sparse coding
    - Independent component analysis
    - Predictive coding
- Conclusion: the brain *may* be trying to find faithful and efficient representations of the natural environment

### 1.4 The Personality of Neurons

#### *Essentially neurobio 101*

- Main character: **cortical neuron**
  - Very small about 25 micron
- **Visual cortex**
  - **Axons form the pyramidal track in motor system**
- **Neuron doctrine**
  - Neuron is fundamental structural and functional unit
  - Neurons are discrete cells
  - Information flows from dendrites to the axon via cell body
- Dendrites are like the inputs
- EPSP – excitatory post-synaptic potential
- A bunch of these get fed into the dendrites and then essentially the summation of these is the action potential
- If some threshold is reached, then we have this *action potential* which is the output
- **Def neuron**
  - Leaky bag of charged liquid
  - Neuron insides enclosed within cell membrane
    - Cell membrane is a lipid bilayer
    - Impermeable to charged ion species
    - BUT there are ionic channels
      - The ionic channels let ions flow in and out
  - Maintains a potential difference across membrane

- Concentration of ion difference leads to -70 mV
- Ionic channels
  - Voltage-gated: prob of opening depends on membrane voltage
  - Chemically-gated: binding to a chemical causes channel to open
  - Mechanically-gated: sensitive to pressure or stretch
- Synapses
  - Junctions between neurons
  - Changes in local membrane potential
- Voltage gated channels cause action potentials
  - Depolarization opens sodium channels
  - Really about the sodium and potassium balance
  - Downward spike of action potential is from the sodium channels



- The wrapping of part of the axons is called *myelin sheath*
- The myelination of axons allows for *fast long-range spike communication*
- **Action potential hops from one non-myelinated region to the next**
  - These non-myelinated regions are called *node of Ranvier*
  - This is essentially active wire → lossless signal propagation

### 1.5 Making Connections: Synapses

- Synapse – connection between two neurons
  - Electrical synapses – gap junctions
    - Helpful for when you need to synchronize
    - Neurons fire simultaneously
  - Chemical synapses – neurotransmitters

- Basis for learning and memory
- Changes the way the other neuron is affected simply by changing density
- Can be excitatory or inhibitory
  - **Def excitatory**
    - Tends to increase the post synaptic membrane potential
    - Tends to excite membrane b
    - Neurotransmitters could be: glutamate
  - **Def inhibitory**
    - Tends to decrease the post synaptic membrane potential
  - So there is a spike, release of neurotransmitter, ion channels open, sodium influx, depolarization
- *Synapses are the basis for memory and learning*
- Allow for learning through: *synaptic plasticity*
  - Hebbian Plasticity
    - If a neuron repeatedly takes part in firing another neuron, then the synapse between those neurons is strengthened
    - *“Neurons that fire together, wire together!”*
    - Evidence: long term potentiation (LTP)
      - Experimentally observed *increase* in synaptic strength
    - Long term depression (LTD)
      - Experimentally observed *decrease* in synaptic strength
    - LTD and LTP are generally confirmed with decrease in EPSP size
  - Synaptic plasticity depends on spike timing!
  - If input is **after** output → LTD
  - If input is **before** output → LTP

### 1.6 Time to Network: Brain Areas and their Function

- Mainly two types of nervous systems
- **Peripheral Nervous System (PNS)**
  - Two main components
  - **Somatic** – nerves connecting to voluntary skeletal muscles and sensory receptors
  - Ex. Moving your arm and hand to shake a friends hand → utilized the SOMATIC nervous system
    - *Afferent Nerve Fibers (incoming)*
      - Axons that carry info away from the periphery to the CNS (central nervous system)
    - *Efferent Nerve Fibers*
      - Carry info from CNS to periphery
  - **Autonomic**
    - Nerves that connect to heart, blood vessels, etc.
    - Guilty of “fight or flight” reaction
- **Central Nervous System (CNS)**

- Spinal Cord + Brain
- **Spinal Cord**
  - Local feedback loops → reflex arc
    - Ex: jumping up when you step on a nail
    - Or jerking at a hot surface
  - Descending motor control signals → activate spinal motor neurons
    - Ex: brain tells your body to walk. Your spinal neurons are the ones that control this. So this way you can walk and also talk.
  - Ascending sensory axons
    - Convey *sensory* information from muscles and skin to the brain
- **BRAIN**
  - Region
  - **Hindbrain – Medulla oblongata, pons, cerebellum**
    - *Medulla oblongata*
      - Breathing, muscle ton
    - *Pons*
      - Connected to cerebellum
      - Involved in sleep and arousal
    - *Cerebellum*
      - EQUILLIBRIUM
      - Language and attention
      - Coordination and timing of voluntary movements
  - **Midbrain and Retic Formation**
    - *Midbrain*
      - Eye movements, visual and auditory reflexes
    - *Reticular Formation*
      - Modulates muscle reflexes
      - Regulates sleep
      - Wakefulness and arousal
  - (near center) **Thalamus and Hypothalamus**
    - *Thalamus*
      - “relay station” for all sensory information to the cortex
      - regulates sleep and wakefulness
    - *Hypothalamus*
      - Right below the thalamus
      - BASIC NEEDS (the four f’s) <- lol:
        - **FIGHTING**
        - **FLEEING**
        - **FEEDING**
        - **MATING**
  - **Cerebrum**

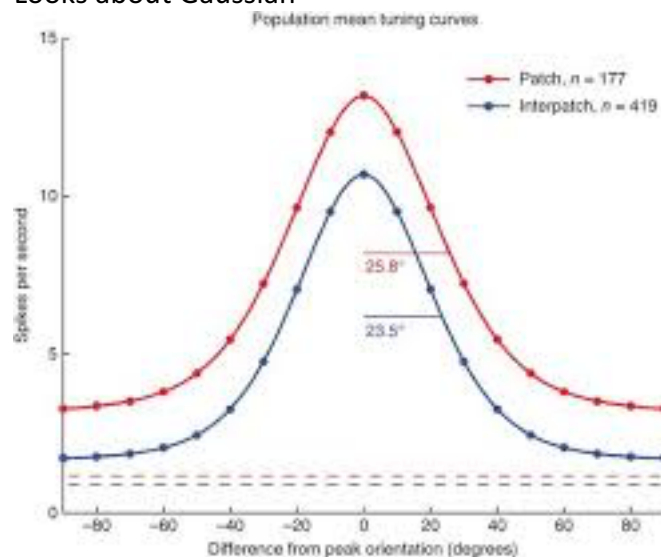
- Consists of cerebral cortex, basal ganglia, hippocampus, and amygdala
- Perception, motor control, cognitive functions, emotions, memory and learning
- *Cerebral Cortex*
  - Layered sheet of neurons
  - 1/8<sup>th</sup> of an inch thick
  - 30 billion neurons. 10,000 synapses each.
  - 300 trillion connections in total
  - Six layers of neurons
- Neural vs Digital Computing
  - The brain is **massively parallelized**
  - **Adaptive connectivity**
  - Digital computing:
    - More sequential via CPUs with fixed connectivity
  - Large computational analogs
    - Information storage: physical/chemical structure of neurons and synapses
    - Information transmission: electrical/chemical signaling
    - Primary computing elements: neurons
    - Computational basis: unknown

## **WEEK 2 – Neural Encoding and Decoding**

### *2.1 What is the Neural Code?*

- Tool for recording from the brain: fMRI
  - Functional magnetic resonance imaging
  - Measures spatial perturbations in the magnetic field
    - The changes are caused from blood oxygenation
    - As blood flows around you can see the underlying neural activity
- EEG's also just show activity for a bunch of neurons
- Calcium imaging is another way to read the neural code
- **What is the actual neural code?**
  - Let's look at the retina
  - *Retina* – sheet of cells at the back of the eyeball
    - Take light from the lens and converts to electrical signals
  - Raster plot – way of visualizing multiple iterations
  - Each neuron encodes a bit of the movie (from the experiment)
- Two questions:
  - Encoding: how does a stimulus cause a pattern of responses?
    - Stimulus → response
    - $P(\text{response} | \text{stimulus})$  *encoding*
  - Decoding: how do the responses tell us about stimulus?

- Response  $\rightarrow$  stimulus
- $P(\text{stimulus} \mid \text{response})$  *decoding*
- Neuron response is some type of average firing rate of generating a spike
- Tuning curve
  - Frequency vs orientation of light
  - Looks about Gaussian



- There is higher order of spatial recognitions
- MRI's highlight different regions when shown faces vs houses
- Tuning curves can be difficult to record
- Building up complex selectivity
  - Brain areas build up the complexity of stimulus representation
  - Geometric in retina and thalamus, to V1 (orientated edges) and then V4.
  - Higher order areas are less sensitive to details such as color or location.
  - This is the idea behind hierarchical features in a feed forward way

## 2.1 Neural Encoding: Simple Models

- Basic coding model
  - linear response
    - $r(t) = \theta * s(t)$  (maybe  $-\theta * s(t - \tau)$ )
    - just going to be delayed and scaled by a little bit
  - Temporal filtering (convolution)
    - We expect response to depend on the **combination of recent inputs**
    - $r(t) = \sum_{k=0}^n s_{t-k} f_k$
    - this is like convolution
    - in fact exact definition. See Cheever's page for refresher.
    - Example:
      - Running average
      - Leaky average

- Spatial filtering
  - Connected with receptive fields
  - So  $r(t) = \sum s_{t-k} f_k$  temporal
  - $r(x, y) = \sum_{x'=-n, y'=-n}^n s_{x-x', y-y'} f_{x', y'}$
  - The receptive field is  $f$ . How similar is it to the receptive field is expressed by  $f$
  - Often our receptive field  $f$ , is going to be a difference of Gaussians
  - Difference of Gaussians really just picks up the edges
- Spatiotemporal filtering
  - Both space and time are going to be best
  - We need a combination
- Another solution is to have a linear filter and a nonlinearity
  - Something like:
  - $r(t) = g(\int s(t - \tau) f(\tau) d\tau)$
  - How do you find the components of the model?

### 2.3 Neural Encoding: Feature Selection

- A good basic coding model: combination of a linear filter and a nonlinear input-output function
- One problem is of dimensionality
- Need to find the feature that drives the neuron
- Just enough so we can learn what really drives cell
- Start with  $s(t)$  and discretize
- What is the right stimulus to use?
  - Gaussian white noise
  - We choose a new Gaussian number at each frequency
  - The prior distribution is the distribution of the stimulus
  - Multivariate Gaussian – Gaussian no matter how we look at it
- Determining linear features -> one good way is to take the average
  - The vector through this average → spike triggered average
  - Then we can project all of the other points and project along that axis
- Linear filtering = convolution = projection
- Looking for stimulus feature  $f$  which is a vector in high dimensional stimulus space
- **Summary: find a feature by:**
  - **Stimulate with white noise**
  - **Reverse correlation to compute spike triggered average**
  - **This is good approximation to our feature**
- Still though how do we compute input / output w.r.t. feature
- $P(\text{spike} \mid \text{stimulus}) \rightarrow P(\text{spike} \mid \text{component of the stimulus extracted by linear filter})$
- Then use Bayes Rule
- $P(\text{spike} \mid s_1) = P(s_1 \mid \text{spike}) P(\text{spike}) / P(s_1)$




- Denominator is called the prior remember
- And  $P(s1 | \text{spike})$  is the spike conditional distribution
- $P(\text{spike})$  is independent of the stimulus
- $P(\text{spike} | s1) = P(s1 | \text{spike}) P(\text{spike}) / P(s1)$ 
  - Let's assume random  $\rightarrow$  this means that is if the blue and red don't change
  - Then we might have filtered out the right feature
  - What we want to see is a nice difference between the prior and the spike conditional
  - This means that our input/output curve will be interesting and we can predict high firing rates
- Let's add the possibility of multiple features
- This essentially means there are several filters
- We could use PCA!! Ahhh
  - This way we get like the main dimensionality
  - As the video puts it, general, famous, and kind of magical tool for discovering low dimensional structure
  - The components correspond to orthogonal set of vectors that span the cloud
  - The important dimensions are some unknown linear combination of dimensions
  - Gives a new basis set to represent the data  $\rightarrow$  lots of compression
  - Here, it is going to be some basis of our features
  - Tangent: eigenfaces!!
    - We can rep almost any new faces as sums of different eigenfaces
- PCA picks out the dimension with the largest amount of variance
- Then we project the rest of the data into the feature space
- We're trying to find interesting features in the retina
- We find an "on" and an "off" feature
- Using this technique, we can plot our data in the two feature axes and we can find the on and the off features
- NOTE: the two features are not the on and off feature themselves, but they allow a coordinate system where we can see the structure

## 2.4 Neural Encoding: Variability

- Recall the Gaussian function:
- $P(x) = Ae^{-\frac{(x-x_0)^2}{2\sigma^2}}$
- When we use something like PCA, making sure that we have a stimulus that's as symmetric as possible with respect to coordinate transformations
- But what if we don't use PCA, and we just look at the prior and the conditional distribution and say, can I find a filter? Meaning, like when I project the stimulus onto it are the distribution and prior as different as possible
  - Standard for measuring the difference between two probability distributions:
  - **KULLBACK-LEIBLER DIVERGENCE (DKL)**

- $D_{KL}(P(s), Q(s)) = \int ds P(s) \log_2 \frac{P(s)}{Q(s)}$
  - So we just want to maximize this
  - Kind of turns into an optimization problem
- Maximally informative dimensions
  - Choose filter to maximize DKL between spike conditional and prior distributions
  - So we just vary our filter around, to maximize the DKL
  - Trying to find a stimulus component that is *as informative* as possible
  - This is a really powerful technique because it can generate
  - HOWEVER, A DOWNSIDE IS THAT THIS IS A VERY TOUGH OPTIMIZATION PROBLEM AND GLOBAL OPTIMIZATION IS TRICKY
- Finding relevant features
  - Single filter determined by conditional average
  - Family of filters from PCA
  - Information theoretic methods that use whole distribution
- Assumption that we make is that every spike is independent of others
  - Bernoulli trials
  - So kind of like a coin flipping
  - Dividing time sample into multiple time bins
  - Sequence of n time bins where  $n = T / \Delta t$
  - **Binomial distribution**
    - $P =$  probability of firing
    - Distribution:  $P_n[k] = p^k (1 - p)^{n-k} \binom{n}{k}$  ( $n \text{ choose } k$ )
    - The  $n \text{ choose } k$  is because we don't care about the way we're arranging those  $k$  spikes
    - Average:  $np$  or  $rT$
    - Variance:  $np(1-p)$
    - Fano factor:  $F = 1$
    - Interval distribution =  $P(T) = r \exp(-rT)$
    - Fano factor – tests if something is a Poisson distribution or not
    - If fano factor == 1: it is Poisson
    - Here, we have defined  $r$  as the rate or probability of per units of time
    - $T$  is our time
    - We do some calculations from binomial and binomial  $\rightarrow$  Poisson
    - Ex problem:
      - Suppose that while a stimulus is present, a neuron's mean firing rate is  $r=4$  spikes / second. If this neuron's spiking is characterized by Poisson spiking, then the prob that the neuron fires  $k$  spikes in  $T$  seconds is given by:
      - $$p(k) = \frac{(rT)^k e^{-rT}}{k!}$$

- What is the prob that when this stimulus is shown for one second the neuron does not fire any spikes?
- $e^{-4}$  bc  $p(0) = 1 * e^{-4} / 1$ 
  - Intervals between spikes have exponential distribution
- Two strong traits of Poisson:
  - Fano factor == 1
  - Interval distribution: exponential distribution of times
- So then we can look at the slope of the number of spikes vs the mean count and then we can look at the slope
- If distributed Poisson, then the slopes should all be 1. So looking at the variance vs the mean count should have a slope of about 1
- Poisson nature of firing and randomness that we need takes care of random background noise
- Poisson assumes spike time independent
- Real neurons have refractory period that prevents the cell from spiking immediately
- Generalized linear model:



- Exponential non linearity → able to find all parameters of the model, using an optimization scheme that is globally convergent
- More generality but model now more complete in another way
- GLM = general linear model
- *Time rescaling theorem*
  - Use Poisson nature to test whether we have captured everything
  - We can predict our output spike intervals and scale them by firing rate that's predicted

- Take interval times and scale them by firing rate
- These new scaled intervals should be distributed like a pure Poisson process
- As a single clean exponential

## QUIZ 2

1. A cosine function is not a linear filtering system
2. **The definition of a spike triggered average for a neuron is The set of stimuli preceding a spike, each averaged over time.**
  - a. **I got this wrong. The correct answer is The averaged stimulus values over a given time before a spike that elicit a spike. That should have been obvious from the python script but alas...**
3. Sampling rate is 1sample/500s. so in 1s / 500 Hz = 0.002 sample period. Sampling period is the inverse of the sampling frequency. This is 2 ms.
4. # time steps in our average vector is 300 ms / width between interval = 300 ms / 2 ms from #3 = 150.
5. Just `len(num_spikes) = 53583`
6. See corresponding code
7. Leaky integration? Because we can see that things are decaying away prior to the spike
8. We can kind of think of this neuron like a capacitor. I had to look this one up because I wasn't sure. But yeah so it's kind of charging up right? So like the best thing is going to be a constant positive value because then it will gradually charge up and fire it's neuron.
9. PCA is the best of the ways

## WEEK 3 – Extracting Information from Neurons: Neural Decoding


### 3.1 Neural Decoding and Signal Detection Theory

- Really going to choose between two cases:
  - Single neuron
  - Range of choices, where there are a few neurons that might be affected by the stimulus
- Also how do we decode in real time
- Famous experiment to determine how noisy sensory information was interpreted
  - Monkey would focus on a screen
  - Watch a pattern of random dots move across the screen
  - Monkey trained → follow the dots. Tracking the dot patterns.
  - Dot pattern is noisy. Hard to tell which way it's going.
  - Fraction that the monkey actually gets right is a function of the coherence. It looks like a sigmoidal function almost
- Signal Detection Theory
  - We can generate some graphs
  - R is the number of spikes in a single trial
  - Two probability distributions dist. Normal
  - $P(r | -)$  and  $P(r | +)$

- We want to map some range of  $r$
- This means some threshold
- The intersection between the two Gaussians would maximize the percentage correct
- $P_{\text{corr}} = P(+ | r \geq z) + p(-) (1 - p(r \geq z | -))$
- False alarms:  $P(r \geq z | -)$
- Good calls:  $P(r \geq z | +)$
- These probabilities  $p(r | -)$  and  $p(r | +)$  are known as the likelihoods
- Choosing the maximum likelihood
- Likelihood ratio
  - Putting a threshold on the likelihood ratio
  - $\frac{p(r|+)}{p(r|-)} > 1$  whenever we choose plus
  - This is the most efficient statistic to use, it has the most power for its size
  - This is called the Neyman-Pearson Lemma
  - [https://en.wikipedia.org/wiki/Neyman%E2%80%93Pearson\\_lemma](https://en.wikipedia.org/wiki/Neyman%E2%80%93Pearson_lemma)
  - Really cool lemma actually
- Seems to be a close correspondence between decoded neural response and monkey's behavior
- So why do we have so many neurons? Tbd
- Log odds!! Ah Zucker talked about this in mobile
- So we have
- $l(s) = \frac{p(s|tiger)}{p(s|breeze)}$
- $\log(l(s)) = \log p(s|tiger) + \log p(s|breeze)$


## Let's just consider for a moment

Now let's say we don't have to decide immediately...



$$l(s) = \frac{P(s|\text{tiger})}{P(s|\text{branch})} \} < 1$$

$$\log l(s) < 0$$



- 3.1 Neural Decoding and Signal Detection Theory
- Firing rates ramp up until a certain sure decision
- But back to our trial...
- What is the actual probability that we have a tiger? It's really low! We need to take into account the *priors*.
- The wind or a tiger?
  - Rods in your eyes can response to light and even a single photon
  - So if we adjust our probability distributions then we can pick out instances when there is a significant difference in firing rate
  - Building in cost
  - We have multiple loss functions
- Loss Functions
  - Loss\_minus = L\_minus P[+|r]
  - Loss\_plus = L\_plus P[-|r]
  - Cut your losses: answer plus when loss\_plus < loss\_minus
  - New criterion for the likelihood ratio:
    - $\frac{p[r|+]}{p[r|-]} > \frac{L_+ P[-]}{L_- P[+]}$

## 3.2 Population Coding and Bayesian Estimation (kind of a tough one to get through)

- Crickets are sensitive to wind. Like wicked sensitive.
- All because of cricket cercal cells.

- These neurons respond with peaks in one of the four cardinal directions, which is 45° to the animal. Left and right, front and back.
- The curves are approx. cosine, so that neurons respond to cosine of angle. Neuron's firing rate is proportional to the projection of the wind velocity.
- Bayesian Inference
  - $p[s|r] = \frac{p[r|s]p[s]}{p[r]}$
  - *a posteriori distribution* = *likelihood function* \* *prior*  $\frac{\text{distribution}}{\text{marginal distribution}}$
  - Maximum likelihood  $s^*$  which maximizes  $p[r|s]$
- Decoding an arbitrary continuous stimulus
  - Assume independence
  - Assume Poisson firing
    - Spikes are random and independent
    - $P_T[k] = \frac{(rT)^k \exp(-rT)}{k!}$
    - Then we want,  $r_a$  to stimulus  $s$
    - That is the firing rate to a stimulus
    - $P(r_a|s) = \frac{(f_a(s)T)^{r_aT} \exp(-f_a(s)T)}{(r_aT)!}$
    - $P(r_a|s) = \prod \frac{(f_a(s)T)^{r_aT} \exp(-f_a(s)T)}{(r_aT)!}$  because we're assuming independence and then we go from  $a = 1$  to  $N$
    - We can take the log
    - The math got pretty hairy so here are some photos

## Maximum likelihood

$$P[r|s] = \prod_{a=1}^N \frac{(f_a(s)T)^{r_a T}}{(r_a T)!} \exp(-f_a(s)T)$$

$$\ln P(r|s) = \sum_{a=1}^N \left\{ r_a T \ln(f_a(s)T) - \underline{f_a(s)T} - \underline{\ln(r_a T)!} \right\}$$

$$\frac{\partial}{\partial s} \ln P(r|s) = T \sum_{a=1}^N r_a \frac{f'_a(s)T}{f_a(s)} \quad \sum_{a=1}^N f_a(s)T = c$$

$$= T \sum_a r_a \frac{f'_a(s)}{f_a(s)} = 0$$



## Maximum likelihood

$$\sum_{a=1}^N r_a \frac{f'(s^*)}{f(s^*)} = 0 \quad \leftarrow \quad \begin{aligned} f_a(s) &= A e^{-\frac{1}{2\sigma_a^2} (s-s_a)^2} \\ f'_a(s) &= A \frac{(s-s_a)}{\sigma_a^2} \cdot e^{-\frac{1}{2\sigma_a^2} (s-s_a)^2} \end{aligned}$$

$$\sum_{a=1}^N \frac{r_a (s-s_a)}{\sigma_a^2} = 0$$

From Gaussianity of tuning curves,

$$s^* = \frac{\sum r_a s_a / \sigma_a^2}{\sum r_a / \sigma_a^2}$$

If all  $\sigma$  are the same

$$s^* = \frac{\sum r_a s_a}{\sum r_a}$$

## Maximum *a posteriori*

Maximize  $\ln p[s|\mathbf{r}]$  with respect to  $s$

$$\ln p[s|\mathbf{r}] = \ln P[\mathbf{r}|s] + \ln p[s] - \ln P[\mathbf{r}]$$


$$\ln p[s|\mathbf{r}] = T \sum_{a=1}^N r_a \ln(f_a(s)) + \ln p[s] + \dots$$

Set derivative to zero, use sum = constant

$$\sum_{a=1}^N r_a \frac{f'(s^*)}{f(s^*)} + \frac{p'(s)}{p(s)} = 0$$

From Gaussianity of tuning curves,

$$s^* = \frac{T \sum r_a s_a / \sigma_a^2 + s_{\text{prior}} / \sigma_{\text{prior}}^2}{T \sum r_a / \sigma_a^2 + 1 / \sigma_{\text{prior}}^2}$$



- And then we want to take the derivative and set that equal to zero to find the most likely value

- Ok I didn't want to write all the equations out in a word doc so here are the pictures
- This method takes care of weighting them based on the variance
- Limitations
  - Tuning curve / mean firing rate
  - Correlations

### 3.3 Reading Minds: Stimulus Reconstruction

... should go back and rewatch

- One day – play back our dreams?
- Extend model to handle varying continuously in time.
- We want to find estimator  $s_{\text{bayes}}$  that gives us best possible
- Introduce error function  $L(s, s_{\text{bayes}})$
- Least squares cost. So just  $L(s, s_{\text{bayes}}) = (s - s_{\text{bayes}})^2$
- Solution:  $s_{\text{bayes}} = \int ds p[s|r] s$
- Reading minds: fMRI
  - Output predicted on BOLD signals (blood oxygen signals)
  - It therefore has a delay
  - ^ that's one way
  - Another way is a motion energy filter

### 3.4 Fred Rieke on Visual Processing in the Retina

- A few rods out of 1000s are contributing signals
- All rods are generating noise
- Averaging would be a disaster
- Have access to rod signal and noise properties
- So we see evidence for a nonlinear threshold between rod and rod-bipolar cells
- Vision is working under conditions where the vast majority are generating noise
  - Want to scale the distributions to take into account the prior probability

### QUIZ 3

- Stimulus  $s$ . Can be one of two values  $s_1$  or  $s_2$ . Firing rate response  $r$ . Under stimulus  $s_1$  repose rate is roughly Gaussian  $\sim N(5, .5^2)$ .  $S_2 \sim N(7, 1)$ .
- It is twice as bad to mistakenly think that it is  $s_2$  rather than  $s_1$ .
  - So this is saying something about where we're thresholding.
- "The disease is very rare. The prior probability of being positive for the disease is therefore very low. MAP (maximum aPRIORI) takes this into account; MLE does not. The mathematics differ in that MAP includes a term for the prior."
  - From my stats.stackexchange question I asked about

**WEEK 4 – Information Theory and Neural Coding****4.1 Information and Entropy**

- Going to start by talking about entropy and information
- How to compute information for neural spike trains
- And what can this tell us about coding
- Ok so back to the monkey example:
  - Information quantifies surprise
  - Some overall prob  $p$  that there's a spike
  - $P(1) = p$
  - $P(0) = 1-p$
  - $\text{Information}(1) = -\log_2 p$
  - $\text{Information}(0) = -\log_2 (1-p)$
- Why does the information have this form?
- Each *bit* of information specifies location by factor of 2
- What we're really doing is multiplying the probabilities
- **Entropy – average information of a random variable**
  - **Measures variability**
  - Units are in bits
  - Entropy counts the yes / no questions
  - $\text{Entropy} = -\sum p_i \log_2 p_i$
  - Or in continuous  $-\int dx p(x) \log_2 p(x)$
- This is essentially just hunting for the binary search
- $H = -\sum p_i \log p_i$
- $p_i = \frac{1}{8}$
- $H = -\sum_{i=1 \text{ to } 8} \frac{1}{8} \log \left(\frac{1}{8}\right)$
- $\sum \frac{1}{8} * -3 = 3$
- Three questions to find car (in example) and that's exactly the entropy
- Maximize the entropy
  - Compute the entropy as a function of the probability  $p$
  - What does having a large entropy do for a code?
  - Gives the most possibility for representing inputs
  - You want to find the value of  $p$  such that  $H$  has a max
  - If  $p = \frac{1}{2}$  then those two symbols are used equally as often
- Entropy tells about intrinsic variability of the outputs
- Week 2 was asking how do we know what our stimulus was
- But now, we need to incorporate our error chances
- Assume the same error
- How much of the entropy is accounted for by these errors?
- Total entropy:  $H[R] = -P(r_+) \log P(r_+) - P(r_-) \log P(r_-)$
- Noise entropy:  $H[R|+] = -q \log q - (1-q) \log (1-q)$

- These stimulus driven entropies are called *noise entropies*
- Amount of entropy that is used in coding the stimulus
- $MI(S, R) = \text{Total entropy} - \text{average noise entropy}$
- $MI = -\sum_r p(r) \log p(r) - \sum_s p(s) [-\sum_r p(r|s) \log p(r|s)]$
- Entropy and information
  - Fixing p
  - Vary the noise probability
  - *When there is no error, the mutual information is 1 to 1. Information is just the entropy of the response.*
  - *As the error rate increases, error probability grows larger and larger.*
  - ***If  $p(r|s) = p(r)$ , the mutual information MI of r and s is zero, because this is saying r and s are independent and therefore no information is gained***
  - ***If response is perfectly predicted, then the MI is 1, because total information is conveyed by 1.***
- Mutual information measures relationship
  - The information quantifies *how independent* R and S.
  - Going to use the Kullback Leibler divergence.
    - This is a measure of the difference between two prob distributions
    - Normally it is between a “true” distribution and a theoretical distribution
    - [https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler\\_divergence](https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler_divergence)
  - $D_{KL}(P, Q) = \int dx P(x) \log \frac{P(x)}{Q(x)}$
  - Going to generalize so that distributions are functions of s and r. So we would need to integrate over both s and r
  - $\int ds dr P(s, r) \log \frac{P(r,s)}{P(r)P(s)} = \int ds dr P(s, r) \log \frac{P(r|s)P(s)}{P(r)P(s)}$
  - $= \int ds dr P(s, r) [\log P(r|s) - \log P(r)]$
  - $= -\int ds dr P(s, r) \log P(r) + \int ds dr P(s) P(r|s) \log P(r|s)$
  - The first bit we can just integrate over s
  - The second term is going to be the entropy of  $P(r|s)$
  - This gives us exactly what is expected
  - $I(S, R) = H[R] - \sum_s P(s) H[R|s]$
- Calculating mutual information
- Take one stimulus s and repeat many times  $\rightarrow$  obtain  $P(R|s)$
- Compute variability due to noise: noise entropy  $\rightarrow H[R|s]$
- Repeat for all s and average  $\rightarrow \sum_s P(s) H[R|s]$
- Compute  $P(R) = \sum_s P(s) P(R|S)$  and theoretical entropy

#### 4.2 Calculating Information in Spike Trains

- Two methods: single spikes, vs lots of spikes
- Mutual information = diff( total response entropy , mean noise entropy)
- Methodology:

- Divide up voltage train into letter size  $\Delta t$  and length  $T$
- Essentially then just have a 1 if we have a spike 0 if not
- From this, compute  $p(w_i)$
- $H[w] = -\sum p(w_i) \log p(w_i)$
- How to sample  $P(S) \rightarrow$  average over time
- For each time, we're going to given set of words  $P(w|s(t))$
- Then we have an average entropy
- Choose length of repeated sample long enough so that the sample the noise adequately
- Information in single spikes is similar to what we just saw in the previous lecture
- After a bit of math and some assumptions, the *information per spike* is:
- $I(r, s) = \frac{1}{T} \int_0^T dt \frac{r(t)}{r_{bar}} \log \frac{r(t)}{r_{bar}}$
- No explicit stimulus dependence (NO NEED FOR CODING / DECODING MODEL)
- The rate  $r$  does not have to mean rate of spikes  $\rightarrow$  can be rate of any event
- Limitations of information:
  - Spike precision, blurs  $r(t)$
  - Mean spike rate

#### 4.3 Coding Principles

- Natural stimuli
  - Huge dynamic range
  - Power law scaling
- Efficient coding:
  - In order to have max entropy output, a good encoder should match its outputs to the distribution of its inputs
  - Should be able to stretch its input axis (IN REALTIME) so that it can accommodate the variations in the overall scaling
- Feature adaption
  - Power spectrum and signal to noise ratio are large factors for the predicted receptive field at certain light levels.
  - Center becomes broader in low light levels
  - Choose filter to maximize Kullback Leibler divergence between spike conditional and prior distributions
- Redundancy reduction
  - Neural systems should be trying to encode as efficiently as possible
  - Maximize the entropy should take into effect the marginal added together
  - Correlations can be good  $\rightarrow$  error correction + correlations help discrimination
- Neuron populations should be as **SPARSE** as possible
  - Let's say we right down a set of basis functions,  $\phi_i$ ,
  - Any image can be expressed as a weighted sum
  - $I(x_{bar}) = \sum_i a_i \phi_i(x) + \epsilon(x)$

- Want to penalize having too many images
- Fourier basis represents in sin and cosine... but not necessarily sparse because the power spectrum is broad
- Sparse code – excites a minimum number of image
- Classic and State of the Art Methods:
  - Models for how stimuli are coded in spikes
  - Models for decoding stimulus from neural
  - Information theory
  - A very quick glance at how coding strategies might shape other things

## **WEEK 5 – Computing in Carbon**

### *5.1 – Modeling Neurons*

- About to delve into circuit diagrams
- Differential equations (largely first order)
- Hodgkin Huxley model
  - Should be a review from biomedical signals
- Basic review of circuit diagrams
- Membrane patch
  - We have a lipid bilayer
    - Like a capacitor
  - Pores
  - Channel
- Cell battery
  - Outside the cell: higher sodium, chlorine and calcium contents
  - Inside higher potassium levels
  - Concentration gradient = battery
    - Nernst Equation  $E = \frac{k_b T}{z q} \ln \frac{[inside]}{[outside]}$
- Currents flow through ion channel

### *5.2 – Spikes*

- What makes a neuron compute?
- Neuron responds to steps and thresholds
  - Uncover the non-linearity
- Gate has subunits that need to be open for things to go through
- *Gating* depends on subunit state
  - $P_k = n^4$
  - $n$  is open prob
  - $1-n$  is closed
- Review of biomedical signals
- Independent probability of being open
- Hodgkin and Huxley's nobel equation

- Specifies conductance for different channels
- Time constant dictates how rapidly each variable corresponds to voltage change
- Hodgkin-Huxley Model
  - Two different places
  - Biophysical realm → ion channel physics, additional channels
  - Simplified models → fundamental dynamics, analytical tractability

### 5.3 – Simplified Model Neurons

- Can one build a large model with lots of neurons
- Capturing the basic
  - Force the model to be linear
  - $dV/dt = f(V) + I(t)$  # nonlinear because of  $f(V)$
  - $dV/dt = -a(V - V_0) + I(t)$
  - Like a passive membrane
  - $C_m dV/dt = -g_L(V - V_e)$
  - Integrated firing model ^
- Exponential integrate-and-fire neuron
- The theta neuron
  - Great for periodic neurons
  - One dimensional
  - $\frac{d\theta}{dt} = 1 - \cos \theta + (1 + \cos \theta) I(t)$
- Two dimensional models
  - Need a phase plane diagram
  - Can find the nullclines – the place where the derivative is equal to zero
  - Fixed point is going to be the intersection between two nullclines
- Various neurons have different firing rates and oscillations

### 5.4 – A Forest of Dendrites (should review)

- Real neurons are brutal to model
- Inject current at the cell body and record effect in dendrites
- So we're looking at the **soma** to see the response at some input
- Inputs that come in at different parts of the dendrite can have very different effects
- Theoretical basis for dendrite communication
  - PDEs!
  - Linear cables
  - Voltage  $V$  is a function of both  $x$  and  $t$
  - Essentially a bunch of circuits distributed along a cable
  - Now a spatial derivative that has to be taken into effect
  - Essentially, the diffusion equation, but we have an additional  $V_m / r_m$
  - Time constant:  $t_m = r_m c_m$
  - Space constant:  $\lambda = \frac{r_m}{r_i}$

- $R_m$  is membrane resistance
- Function decays rapidly as a function of space
  - Geometry can be extremely complicated → cable equation
  - Ion channels
  - Solution: divide and conquer
  - Each compartment = one  $dV/dt$  equation
  - If branches open a certain branching ratio, can replace each pair of branches with a single cable segment with equivalent surface area and electronic length
- Ion channels introduce the nonlinearity
- Dendrites can add a lot to neuronal computation
  - Logical operations
  - Low pass filter, attenuation
  - Coincidence detection
  - Segregation, amplification
- Example:
  - Delay lines in sound localization

#### *Eric Shea-Brown on Neural Correlations and Synchrony*

- This guy seems good
- Encoding via spikes
- Eye → optic nerve → lateral geniculate nucleus (LGN) → visual cortex
- Tuning curve – firing rates as a function of the angle of some stimulus
- Got a bunch of neurons, have a tuning curve, also variance around that mean
- Two statistics
- Still can quantify similar statistics
- **Pairwise correlation → departure from independence**
  - **Label the spike counts**
  - Pierson correlation coefficient
  - Or just correlation coefficient
  - Then you ask if that number is 0 or non-zero
  - **Correlation can degrade signal encoding**
- Turns out that you can apply this technique to numerous neurons
  - Compute the signal to noise ratio
    - Mean / variance
  - SNR → going to grow with  $M$  (number of neurons)
  - Then also observe the correlation coefficient
- Pairwise couples of entire population

### **WEEK 6 – Computing with Networks**

#### *6.1 Modeling Connections between Neurons*

- Linear filter model of a synapse



- See online notes for this lecture
- Just listened to the audio

## 6.2 Introduction to Network Models

- Learned that neurons use synapses to connect
- Learned how to model with differential equations
- **FEEDFORWARD VS RECURRENT**
- Modelling networks
  - **Spiking Neurons**
    - Pro: Learning based on spike timing
    - Pro: Spike correlations
    - Con: computationally expensive
  - **Firing-rate outputs (real valued outputs)**
    - Greater efficiency, scales well to large networks
    - Ignore spike timing
  - How are they related?
- Synapse **b**
- Input spike train  $\rho_b(t)$
- $\rho_b(t) = \sum_i \delta(t - t_i)$
- $g_b(t) = g_{b,max} \sum_{t_i < t} K(t - t_i) = g_{b,max} \int_{-\infty}^t K(t - \tau) p_b(\tau) d\tau$
- From single synapse to multiple synapses:
  - Each synapse has a synaptic weight
  - Assume no nonlinear interactions
  - Then total synaptic current
  - $I_s(t) = \sum_{b=1}^N w_b \int_{-\infty}^t K(t - \tau) p_b(\tau) d\tau$
  - We go from spike train, to firing rate
  - ***This would fail if there were correlations or synchronies***
- Suppose synaptic filter  $K$  is exponential
- Firing-rate-based network model
- **Output firing rate changes like:**  $\tau_r \frac{dv}{dt} = -v + F(I_s(t))$
- **Input current changes like:**  $\tau_s \frac{dI_s}{dt} = -I_s + w * u$
- Weights matrix **w**
- To get steady state, we need to set both of these equal to zero
- Static input:  $v_{ss} = F(w \cdot u)$
- **THE RICH DYNAMICS THAT ARE ACTUALLY IN THE SYNAPTIC CURRENT ARE REPLACED WITH A SIGMOIDAL FUNCTION FOR ARTIFICIAL NEURAL NETWORKS**
- **THAT'S ONE OF THE BIG DISTINCTIONS**
- **HENCE ARTIFICIAL**
- BIG ASSUMPTIONS THAT THE SYNAPSES ARE RELATIVELY FAST
- Multiple output neurons

- Then we have an input vector and an output vector
- $V$  is now a vector.  $W$  becomes our weight matrix.
- This has all been **FEEDFORWARD NETWORKS**
- $\tau \frac{dv}{dt} = -v + F(Wu + Mv)$
- For feedforward networks,  $M$  is a matrix of zeros
- There's no like passback without recurrent networks!!
- Linear Feedforward Network
  - Steady state:  $v_{ss} = Wu$
- **Edge detectors in the brain**
  - **Primary visual cortex (V1)**
  - **Receptive fields in V1 have edge detection**

### 6.3 The Fascinating World of Recurrent Networks

- Want to find out how the output  $v(t)$  behaves for different  $M$
- Eigen vectors to the rescue!
- $\tau \frac{dv}{dt} = -v + h + Mv$
- Idea use eigenvectors of  $M$  to solve differential equation for  $v$
- Suppose  $N \times N$  matrix  $M$  is symmetric
- If  $m$  is symmetric,  $M$  has  $N$  orthogonal eigenvectors  $e_i$  and  $N$  eigenvalues  $\lambda_i$
- It is useful for them to be orthonormal because then we can write our output vector using eigenvectors
  - $v(t) = \sum_{i=1}^N c_i(t) e_i$
  - Complete expression:  $c_i(t) = h \text{ times } e^{\frac{1}{1-\lambda_i}} (1 - \exp(-\frac{t(1-\lambda_i)}{\tau})) + c_i(0) \exp(-\frac{t(1-\lambda_i)}{\tau})$
  - If any of the  $\lambda_i$  is greater than 1  $\rightarrow$  network explodes
  - If all of them are less than 1, network is stable and  $v(t)$  converges to some steady state value
- Network performs winner-takes-all input selection
- Gain modulation in the nonlinear network
  - Adding a constant amount to the input  $h$  **multiplies** the output
- Memory in nonlinear network
  - Network maintains some short term memory
- Nonsymmetrical recurrent networks
  - Network of excitatory and inhibitory neurons
- Linear stability analysis
  - Stability matrix
  - **THIS IS JUST THE JACOBIAN MATRIX**

**NOTE: could not figure out one on the quiz. Posted on stack.**

<http://stackoverflow.com/questions/41492020/finding-the-steady-state-output-of-a-linear-recurrent-network>

## **WEEK 7 – Networks that Learn: Plasticity in the Brain & Learning**

### *7.1 Synaptic Plasticity, Hebb's Rule, and Statistical Learning*

- Long term potentiation (LTP) – experimentally observed increase in synaptic strength that last for hours or days
- Long term depression (LTD) – experimentally observed decrease in synaptic strength that last for hours or days
- **Hebb's Learning Rule**
  - If neuron A takes part in firing neuron B, then the synapse is *strengthened*
  - Formulation as a mathematical model
    - Let's start with linear feed forward model
    - We have a synaptic weight vector
    - Basic hebb rule
    - $\tau_w \frac{dw}{dt} = u v$
    - Discretization:  $w_{i+1} = w_i + \epsilon * uv$
    - Hebb rule only increases synaptic weights (LTP)
  - Learning rules are NOT stable
  - W grows without bound
  - Covariance rule can both increase and decrease
- Start with the averaged Hebb rule:  $\tau_w \frac{dw}{dt} = Q w$
- Solve this equation to find **w(t)** using eigenvectors
- Substitute in Hebb rule differential equation and simplify as before
- Synaptic weight vector is a linear combination
- Has terms that are exponentially dependent on the values of the correlation matrix
- *For large t, largest eigenvalue term dominates*
- *For Oja's Rule:  $w(t) = \frac{e_1}{\sqrt{\alpha}}$*
- Thus we have shown the brain can do statistics
- Hebbian Learning implements **principal component analysis (PCA)**
- **Hebbian learning learns a weight vector aligned with the principal eigenvector of input correlation/covariance matrix**
  - DIRECTION OF MAXIMUM VARIANCE

### *7.2 Introduction to Unsupervised Learning*

- Can neurons learn to represent clusters
- Feedforward network with two neurons
- Most active neuron in the network
  - The one whose weight vector is closest to an input
  - We can show that by looking at the Euclidean distance between vectors

- Given a new input, we can set the weight vector to the running average of all inputs IN THAT CLUSTER
- Then you pick the most active neuron
- Competitive learning and self organizing maps
  - Also known as Kohonen maps
  - Given an input, pick the winning neuron
  - Update weights for that neuron AND the other neurons in the neighborhood of the winning neuron
    - What do we mean by neighborhood?
    - We have locations assigned and the neighboring ones like literally on a 2d grid
- Unsupervised learning
  - We have causes  $\mathbf{v}$
  - Data points  $\mathbf{u}$
  - You kind of assume that there are multiple gaussians given by some prior
  - Mixture of gaussians model
  - **Goal:** learn a good generative model for the data you are seeing
    - Mimic the data generation process
  - **General approach:**
    - Given data  $\mathbf{u}$ , need to
      - Estimate causes  $\mathbf{v}$
      - Learn parameters  $\mathbf{G}$
- Algorithm for learning the parameters
- Expectation-Maximization algorithm:
  - Iterating through expectation step
  - Then the maximization step
  - **E step – computing the posterior distribution of  $\mathbf{v}$  for each  $\mathbf{u}$** 
    - $p[\mathbf{v}|\mathbf{u}; \mathbf{G}] = \frac{p[\mathbf{u}|\mathbf{v}; \mathbf{G}]p[\mathbf{v}; \mathbf{G}]}{p[\mathbf{u}; \mathbf{G}]}$
    - **soft competition**
  - **M step – charge parameters  $\mathbf{G}$  using results from E**
    - Just updating the mean, variance, and the prior

### 7.3 Sparse Coding and Predictive Coding

- Hebb's learning rule implements principal component analysis
- How do we learn models of natural images?
  - Eigenvectors or principal components
  - Turk and Pentland
  - Eigenvectors of the input covariance matrix
  - Any face image can just be a linear summation of the eigenfaces
  - It's a basis

- This could be great for compression! But not great to extract the local components
  - Edges in a scene
  - Can't get that from an eigenvector analysis
- Define the generative model: likelihood
  - Linear model
  - $u = Gv + \text{noise}$
  - You're generating the likelihood based on a probabilistic model
  - **A lot of machine learning algorithms and things in engineering want to minimize the log of the likelihood**
  - $p[u|v; g] = -\frac{1}{2}|u - Gv|^2 + c$
  - if you MINIMIZE the squared reconstruction error you are MAXIMIZING the likelihood of the data
  - Prior
    - Can make some assumptions
    - Assume the causes  $v_i$  are independent
    - For any input, we want only a few causes  $v_i$  to be active
    - **SPARSE DISTRIBUTION**
      - Also called *super-Gaussian distribution*
      - Very sharp
      - You're taking the exponential of a Gaussian
      - $p[v] = c \cdot \prod \exp(g(v_i))$
  - Bayesian approach to **finding  $v$  and learning  $G$** 
    - Going to maximize the posterior probability of causes
    - Equivalently, maximize the log posterior
    - $F(v, G) = -\frac{1}{2}|u - Gv|^2 + \sum_i g(v_i) + K$ 
      - maximize  $F$  with respect to  $v$ , keeping  $G$  fixed
      - maximize  $F$  with respect to  $G$ , keeping  $v$  from above
      - This is similar to the EM algorithm
      - *Normally, we just use gradient ascent*
      - $\frac{dF}{dv} = G^T(u - Gv) + g'(v)$
      - firing rate dynamics
      - $\tau \frac{dv}{dt} = G^T(u - Gv) + g'(v)$
      - First term is the error,  $Gv$  is the prediction
      - It converges to a stable value
  - Learning the synaptic weights  $G$ 
    - $\tau_G \frac{dG}{dt} = (u - Gv)v^T$
    - This is the Hebbian term
    - This is almost identical to Oja's rule for learning

- **Why isn't this network just doing principal component analysis like Oja's rule?**
- Answer: Network is trying to compute a sparse representation of the image
- Learning G for Natural Images
  - **The basis vectors are a bunch of bars**
  - **Like one hot vectors**
  - **The  $g_i$  look like local edge or bar features SIMILAR TO RECEPTIVE FIELDS IN PRIMARY VISUAL CORTEX**

## WEEK 8 – Learning from Supervision and Rewards

### 8.1 Neurons as Classifiers and Supervised Learning

- The classification problem
- Example: classifying images as faces
  - What is we just group them as +1 and -1. Could we draw a line to separate those groups?
- Recall: the idealized neuron
  - It's essentially thresholding
  - Inputs:  $u_i$ ; synaptic weights:  $w_i$ , and if  $\sum_i w_i u_i > \mu$  then we have an output spike
- **This is called the "perceptron"**
  - We have inputs that are either +1 or -1
  - We can build the equation  $\sum_i w_i u_i - \mu = 0$  which is a hyperplane formula
  - *Perceptrons can classify*
- So the question becomes: how do we learn the weights and the threshold?
- Perceptron learning rule:
  - Adjust  $w_i$  and  $\mu$  according to output error ( $v^d - v$ ):
  - $\Delta w_i = \epsilon(v^d - v)u_i$  for positive input; increases weight if error is positive decreases weight if error is negative
  - $\Delta \mu = -\epsilon(v^d - v)$  decreases threshold if error is positive and increases if error is negative
- Great! So perceptrons learn any functions?
- Let's think about XOR:
  - Can't really do it
  - There's no line we can draw
  - **Perceptrons can only classify linearly separable data**
- However, **we can use multilayer perceptrons**
- What about continuous outputs?
  - Sigmoid functions!
  - Output:  $v = g(w^T u) = g(\sum_i w_i u_i)$
  - Sigmoid output function:  $g(a) = \frac{1}{1 + e^{-\beta a}}$

- Range of  $-\infty$  to infinity and then compresses everything to one
- Beta controls the slope
- Learning Multilayer sigmoid networks
  - You could learn weights that minimize the output error
  - $E(W, w) = \frac{1}{2} \sum_i (d_i - v_i)^2$
  - Use gradient descent!!
  - How do we change the weights for the hidden layer?
  - **Backpropagation learning rule**
  - $\Delta w_{jk} = -\epsilon \frac{dE}{dw_{jk}}$
  - The answer essentially lies in the chain rule from calculus
  - Example:
    - $\frac{dE}{dw_{jk}} = \frac{dE}{dx_j} \cdot \frac{dx_j}{dw_{jk}}$
  - The error propagates down through the neural network
  - We should see the entire hidden layer affected

## 8.2 Reinforcement Learning: Predicting Rewards

- We learn by trial and error. Rewards are part of this
- We have some state, some reward, and some action
- We need to pick the action that will maximize our future reward
- ***Pavlov and his dog***
  - Classic conditioning experiments
  - Training: bell  $\rightarrow$  food
  - After: bell  $\rightarrow$  salivate
  - But how do we predict rewards delivered *sometime after* the stimulus?
- Want to have some neuron who predicts the expected total future reward?
- **Key idea: utilize dynamic programming**
  - We don't know our future rewards so we need to approximate
  - Learning the weights according to which  $v(t)$  is calculated
  - **Temporal difference (TD) learning rule**
    - $\Delta w(\tau) = \epsilon [r(t) + v(t+1) - v(t)] u(t - \tau)$
    - We have a temporal difference because we have our future prediction and our current prediction

## 8.3 Reinforcement Learning: Time for Action

- How does the brain use reward information to actually select the actions?
- Learn a state-to-action mapping or a **policy**
  - $\pi(u) = a$
  - Should maximize the expected total future reward
  - $< \sum_{\tau=0 \text{ to } T-t} r(t + \tau) >$
- However, that's using a random policy

- Values should act as surrogate immediate rewards → locally optimal choice leads to globally optimal policy
- **Markov Environment**
  - The next state only depends on the current state and the current action
  - This is closely related to dynamic programming
- **Putting it all together: Actor - Critic Learning**
  - Two separate components:
    - Actor (selects action and maintains policy)
    - Critic (maintains value of each state)
  - 1. Critic Learning (Policy Evaluation)
    - Value of state  $u = v(u) = w(u)$
    - $w(u) \leftarrow w(u) + \epsilon[r(u) + v(u') - v(u)]$
  - 2. Actor Learning (Policy Improvement)
    - $P(a; u) = \frac{\exp(\beta Q_a(u))}{\sum_b \exp(\beta Q_b(u))}$
    - probabilistically select an action  $a$  at state  $u$
    - This is like a softmax function
    - It lets us explore all possibilities
  - Then we repeat steps 1 and 2

End of class